

```
#Working with Simple Linear Regression  
#Risk analysis of space shuttle
```

```
import os  
os.chdir("C:/Users/Mohan/Desktop/Datasets")
```

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
%matplotlib inline
```

```
challenger=pd.read_csv('challenger.csv')
```

```
challenger
```

	o_ring_ct	O.ring.failures	temperature	pressure	launch_id
0	6	0	66	50	1
1	6	1	70	50	2



```
#Assign ring failure values to an object
```

```
O_ring_failures=challenger['O.ring.failures']
```

```
O_ring_failures
```

```
0    0
1    1
2    0
3    0
4    0
5    0
6    0
7    0
8    1
9    1
10   1
11   0
12   0
13   2
14   0
15   0
16   0
17   0
18   0
19   0
20   0
21   0
22   1
```

```
Name: O.ring.failures, dtype: int64
```

```
19    6    0    79    200    20
```

```
#Assign temperature values to 'temp' object
```


```
temp=challenger['temperature']
```

```
temp
```

```
0    66
1    70
2    69
3    68
4    67
5    72
6    73
7    70
8    57
9    63
10   70
11   78
12   67
13   53
```

```
14    67
15    75
16    70
17    81
18    76
19    79
20    75
21    76
22    58
Name: temperature, dtype: int64
```

```
import matplotlib.pyplot as plt
%matplotlib inline
plt.plot(temp,O_ring_failures,'o')
plt.ylabel("O_ring_failures")
plt.xlabel("Temperature")
```



```
challenger.corr()
```

	o_ring_ct	O.ring.failures	temperature	pressure	launch_id
o_ring_ct	NaN	NaN	NaN	NaN	NaN
O.ring.failures	NaN	1.000000	-0.725671	0.220326	-0.011993
temperature	NaN	-0.725671	1.000000	0.039818	0.230770
pressure	NaN	0.220326	0.039818	1.000000	0.839932
launch_id	NaN	-0.011993	0.230770	0.839932	1.000000




```
mean_O_ring_failures = challenger['O.ring.failures'].mean()

mean_O_ring_failures

0.30434782608695654
```

```
plt.plot(temp,O_ring_failures,'o')
```

```
plt.ylabel("O_ring_failures")
plt.xlabel("Temperature")
plt.axhline(mean_O_ring_failures, color='r', linestyle='-')
plt.show()
```



```
import statsmodels.api as sm
model=sm.OLS(O_ring_failures,temp).fit()
```

```
#Obtain model summary
model.summary()
```

OLS Regression Results

<b>Dep. Variable:</b>	O.ring.failures	<b>R-squared (uncentered):</b>	0.178
<b>Model:</b>	OLS	<b>Adj. R-squared (uncentered):</b>	0.140
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	4.755
<b>Date:</b>	Mon, 10 Oct 2022	<b>Prob (F-statistic):</b>	0.0402
<b>Time:</b>	06:24:21	<b>Log-Likelihood:</b>	-19.595
<b>No. Observations:</b>	23	<b>AIC:</b>	41.19
<b>Df Residuals:</b>	22	<b>BIC:</b>	42.33
<b>Df Model:</b>	1		

**Covariance Type:** nonrobust

	coef	std err	t	P> t	[0.025	0.975]
<b>temperature</b>	0.0038	0.002	2.181	0.040	0.000	0.007

**Omnibus:** 14.053 **Durbin-Watson:** 1.851

**Prob(Omnibus):** 0.001 **Jarque-Bera (JB):** 12.853

**Skew:** 1.623 **Prob(JB):** 0.00162

**Kurtosis:** 4.696 **Cond. No.** 1.00

Notes:

[1]  $R^2$  is computed without centering (uncentered) since the model does not contain a constant.


[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

## ▼ Observation

As temeparture increases by 1 degree, O.ring.failures increase by 0.0038

```
import seaborn as sns
sns.pairplot(challenger)
```

<seaborn.axisgrid.PairGrid at 0x7fcdd3df6110>




## ▼ Observations

1.The histogram on the diagonal allows us to see the distribution of a single variable 2.The scatter plots on the upper and lower triangles show the relationship (or lack thereof) between two variables.

```
sns.distplot(challenger['temperature'])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning:
  warnings.warn(msg, FutureWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7fcdd3873690>
```




```
# Plot the correlation using heatmap
```

```
corr = challenger.corr()
```

```
sns.heatmap(corr,xticklabels=corr.columns,yticklabels=corr.columns)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fcdd380aed0>
```



## ► Observations

Black color represents negative correlation which exists between temperature and O-ring failures

[ ] ↳ 12 cells hidden

## ▼ Linear Regression with sklearn

```
from sklearn import linear_model as lm
model=lm.LinearRegression()
results=model.fit(X_train,y_train)
```

```
accuracy = model.score(X_train, y_train)
print('Accuracy of the model:', accuracy)
```

```
Accuracy of the model: 0.496692522077835
```

```
#Print coefficients
print('intercept:', model.intercept_)
print('slope:', model.coef_)
```

```
intercept: 3.874282260501661
slope: [-0.05197945]
```

## ▼ Observation

The slope value -0.051 means that the predicted O-ring failures reduces by -0.05 when x(temperature) rises by one degree

```
#Check model score
model
```

```
LinearRegression()
```

```
X_test
```

	temperature	pressure	
20	75	200	
17	81	200	
3	68	50	
13	53	200	
19	79	200	

X\_test.shape

(20, 1)

4 67 50

y\_test.shape

(10,)

X\_test = X\_test.values.reshape((-1,1))

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-97-6fc1b52b26db> in <module>
----> 1 X_test = X_test.values.reshape((-1,1))
```

**AttributeError:** 'numpy.ndarray' object has no attribute 'values'

SEARCH STACK OVERFLOW

#Predictions from the model

predictions = model.predict(X\_test)

print('predicted O.ring.failures:', predictions, sep = '\n')

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-95-111a443fb772> in <module>
      1 #Predictions from the model
----> 2 predictions = model.predict(X_test)
      3 print('predicted O.ring.failures:', predictions, sep = '\n')
```

3 frames

```
/usr/local/lib/python3.7/dist-packages/sklearn/base.py in _check_n_features(self, X,
reset)
    399         if n_features != self.n_features_in_:
    400             raise ValueError(
--> 401                 f"X has {n_features} features, but {self.__class__.__name__}
    "
    402                 f"is expecting {self.n_features_in_} features as input."
    403             )
```

**ValueError:** X has 1 features, but LinearRegression is expecting 2 features as input.

y\_test



```
#Visualize the predictions
plt.scatter(y_test, predictions)
```

## ▼ Observations

A linear model has been obtained

```
#Other way for prediction
y_pred = model.intercept_ + model.coef_ * X_test
print('predicted response:', y_pred, sep='\n')
```

```
predicted response:
[[-0.02417649]
 [-0.33605319]
 [ 0.33967966]
 [ 1.11937141]
 [-0.23209429]
 [ 0.23572076]
 [ 0.23572076]
 [ 0.39165911]
 [ 0.28770021]
 [-0.07615594]]
```

```
#Define new data instance
Xnew = [[30]]
```

```
#Make a Prediction
ynew = model.predict(Xnew)
```

```
#Show the inputs and predicted outputs
print("New Temperature=%s, Predicted O.ring.failures=%s" % (Xnew,ynew))
```

```
☞ New Temperature=[[30]], Predicted O.ring.failures=[2.31489876]
```

```
#Define new data instance
Xnew = [[70]]
```

```
#Make a Prediction
ynew = model.predict(Xnew)
```

```
#Show the inputs and predicted outputs
print("New Temperature=%s, Predicted O.ring.failures=%s" % (Xnew,ynew))
```

```
New Temperature=[[70]], Predicted O.ring.failures=[0.23572076]
```

```
#Evaluating the model
from sklearn.metrics import mean_squared_error, r2_score
X_train = X_train.reshape(-1,1)
y_train_prediction = model.predict(X_train)
```

```

X_test = X_test.reshape(-1,1)
y_test_prediction = model.predict(X_test)

# printing values
print('Slope:' ,model.coef_)
print('Intercept:', model.intercept_)
print("\n")

# model evaluation for training set
import numpy as np
rmse_training = (np.sqrt(mean_squared_error(y_train, y_train_prediction)))
r2_training = r2_score(y_train, y_train_prediction)

print("The model performance for training set")
print("-----")
print('RMSE is {}'.format(rmse_training))
print('R2 score is {}'.format(r2_training))
print("\n")

# model evaluation for testing set
rmse_testing = (np.sqrt(mean_squared_error(y_test, y_test_prediction)))
r2_testing = r2_score(y_test, y_test_prediction)

print("The model performance for testing set")
print("-----")
print('Root mean squared error: ', rmse_testing)
print('R2 score: ', r2_testing)

Slope: [-0.05197945]
Intercept: 3.874282260501661

The model performance for training set
-----
RMSE is 0.32743461522828027
R2 score is 0.496692522077835

The model performance for testing set
-----
Root mean squared error: 0.4404461397642911
R2 score: 0.5268468243091087

# plotting values

# data points
plt.scatter(X, y)
plt.xlabel('Temeperature')
plt.ylabel('O.ring.failures')

```

Text(0, 0.5, 'O.ring.failures')



```
X= X.reshape(-1,1)
y_predicted = model.predict(X)
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-70-4f01f4a3466f> in <module>
----> 1 X= X.reshape(-1,1)
      2 y_predicted = model.predict(X)

/usr/local/lib/python3.7/dist-packages/pandas/core/generic.py in __getattr__(self,
name)
    5485         ):
    5486             return self[name]
-> 5487         return object.__getattribute__(self, name)
    5488
    5489     def __setattr__(self, name: str, value) -> None:
```

**AttributeError:** 'Series' object has no attribute 'reshape'

SEARCH STACK OVERFLOW

```
# predicted values
plt.plot(X, y_predicted, color='r')
plt.show()
```

```
-----
NameError                                    Traceback (most recent call last)
<ipython-input-89-839323038b49> in <module>
      1 # predicted values
----> 2 plt.plot(X, y_predicted, color='r')
      3 plt.show()
```

**NameError:** name 'y\_predicted' is not defined

SEARCH STACK OVERFLOW

## ► Multiple Linear Regression

[ ] ↳ 2 cells hidden

## ► Observation

1.This model has a higher R-squared compared to simple linear model against temperature and O.ring.failure 2.However in this model both temeprature and pressure features became statistcially insignificant to predict O.ring.failure 3.As pressure increases by 1 atmosphere, O.ring.failures increase by 0.0031 and as temperature increases by 1 degree, O.ring.failures decrease by -0.0030

[ ] ↪ 12 cells hidden

## ► Observation

With both temeprature and pressure the model is not linear

[ ] ↪ 2 cells hidden

[Colab paid products](#) - [Cancel contracts here](#)

✓ 11s completed at 7:35 PM

